

Job Description — NLP Developer / NLP Engineer

Location: Remote / Hybrid | Experience: 0–3 years | Employment: Full-time

About the Role

We are looking for an NLP Developer (NLP Engineer) to build, fine-tune, and deploy Natural Language Processing models for real-world applications such as text classification, information extraction, semantic search, and chatbots. You will work closely with product, data, and engineering teams to deliver high-quality NLP features.

Key Responsibilities

- Build NLP pipelines for preprocessing, tokenization, training, evaluation, and inference.
- Develop and fine-tune Transformer-based models (BERT, RoBERTa, DistilBERT, T5, etc.).
- Implement solutions for tasks such as classification, NER, summarization, Q&A;, and semantic similarity.
- Create embedding-based search and retrieval using vector databases (FAISS, Pinecone, Weaviate, Milvus).
- Optimize models for latency and deploy them as APIs (FastAPI/Flask) or microservices.
- Track experiments, evaluate models, and monitor performance in production.
- Work with large-scale text datasets; perform cleaning, normalization, and augmentation.
- Collaborate with frontend/backend teams to integrate NLP features into products.

Required Skills

- Strong Python programming skills.
- Good understanding of NLP fundamentals (tokenization, embeddings, sequence models).
- Hands-on experience with Hugging Face Transformers, PyTorch or TensorFlow.
- Experience with text preprocessing using NLTK / spaCy / regex.
- Familiarity with vector embeddings and semantic search concepts.
- Basic knowledge of REST APIs and deploying models using FastAPI/Flask.
- Working knowledge of Git and software engineering best practices.
- Understanding of evaluation metrics (F1-score, precision/recall, BLEU/ROUGE).

Preferred / Bonus Skills

- Experience with LLMs, prompt engineering, and RAG pipelines.
- Experience with MLflow / Weights & Biases for experiment tracking.
- Familiarity with Docker and basic MLOps workflows.
- Cloud experience (AWS/GCP/Azure) for deployment and storage.
- Knowledge of SQL and data querying.
- Experience with OCR and document understanding pipelines.

Tech Stack (Typical)

- Python, PyTorch, Hugging Face Transformers
- spaCy / NLTK, Pandas, NumPy
- FastAPI, Docker
- Vector DB: FAISS / Pinecone / Weaviate / Milvus
- GitHub, Linux

Interview Process (Example)

- Short screening call
- Python + ML/NLP fundamentals round
- Practical assignment (small NLP task) or live coding
- Final interview with hiring manager

How to Apply

Share your resume, GitHub (if available), and a short note describing an NLP project you have worked on.